

Imports & Notebook Setup

We will import the necessary packages utilized for modeling/evaluation in this notebook. We also make sure that our visualizations are exported to results/figures/modeling & evaluation/ directory, and metrics are exported to results/metrics directory:

```
In [1]:  
import os  
import sys  
from pathlib import Path  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import (  
    accuracy_score,  
    precision_score,  
    recall_score,  
    f1_score,  
    confusion_matrix,  
    classification_report  
)  
from sklearn.preprocessing import LabelEncoder  
  
# Allow imports from ../src directory  
sys.path.append(os.path.abspath("."))  
from src.train_models import (  
    build_logistic_regression_model,  
    build_random_forest_model,  
    build_xgboost  
)  
from src.preprocessing import (  
    train_val_test_split,  
    build_preprocessor  
)  
FIG_DIR = "../results/figures/modeling & evaluation/"  
METRICS_DIR = "../results/metrics/"  
os.makedirs(FIG_DIR, exist_ok=True)  
os.makedirs(METRICS_DIR, exist_ok=True)  
_plt_original_show = plt.show  
_plt_fig_counter = {"count": 0}  
def _save_and_show(*args, **kwargs) -> None:  
    _plt_fig_counter["count"] += 1  
    filename = os.path.join(FIG_DIR, f"figure_{_plt_fig_counter['count']}")  
    plt.savefig(filename, dpi=300, bbox_inches="tight")  
    _plt_original_show(*args, **kwargs)  
plt.show = _save_and_show
```

Compilation complete
Compilation complete

- `pandas` : data wrangling and CSV loading
- `numpy` : array and numerical operations
- `matplotlib & seaborn` : visualizations and plots
- `os, sys, pathlib` : directory management and custom imports
- `sklearn` : train/test splitting, evaluation metrics and label encoding.
- `src.preprocessing` : utilities for feature/target separation, stratified splitting, and building preprocessing pipelines.
- `src.train_models` : Logistic Regression, Random Forest and XGBoost.

Feature and Target Separation

Load `ais_data_model_ready` and separate features (X) and target (y)

```
In [2]: model_df = pd.read_csv("../data/ais_data_model_ready.csv")
TARGET_COL = "navigationalstatus"
X = model_df.drop(columns=[TARGET_COL])
y = model_df[TARGET_COL]
print(f"Number of features: {X.shape[1]}")
print(f"Number of samples: {X.shape[0]}")
print(f"List of features names: {X.columns.tolist()}")
print(f"Target variable: {y.name}")
```

```
Number of features: 7
Number of samples: 325470
List of features names: ['sog', 'cog', 'heading', 'width', 'length', 'draught', 'shiptype']
Target variable: navigationalstatus
```

Train / Validation / Test Split

Two-stage stratified split to preserve class distribution across all splits:

1. Stage 1 (80/20): 80% -> Train+Val, 20% -> Test
2. Stage 2 (80/20 of the 80%): 80% of 80% -> Train (64%), 20% of 80% -> Val (16%)

Final Distribution

- Training: 64%, Validation: 16%, Test: 20%

Function Parameters

- `test_size` : 20% reserved for final evaluation (untouched during development)
- `val_size` : 16% for hyperparameter tuning
- `random_state` : Ensures reproducibility
- Stratified: Preserves imbalanced class distribution across all splits

The function returns 6 variables: `X_train, X_val, X_test, y_train, y_val, y_test`

```
In [3]: X_train, X_val, X_test, y_train, y_val, y_test = train_val_test_split(
    X,
    y,
    test_size=0.2, # 20% reserved for final evaluation (untouched during de-
    val_size=0.16, # 16% for hyperparameter tuning
    random_state=42
)
print("Train/val/test split complete!")
```

Train/val/test split complete!

```
In [4]: # Print split information
print(f"Total samples: {len(X)}")
print(f"\nTrain set:")
print(f"  Shape: {X_train.shape}")
print(f"  Percentage: {len(X_train)/len(X)*100:.1f}%")
print(f"\nValidation set:")
print(f"  Shape: {X_val.shape}")
print(f"  Percentage: {len(X_val)/len(X)*100:.1f}%")
print(f"\nTest set:")
print(f"  Shape: {X_test.shape}")
print(f"  Percentage: {len(X_test)/len(X)*100:.1f}%")
```

Total samples: 325470

Train set:
 Shape: (208300, 7)
 Percentage: 64.0%

Validation set:
 Shape: (52075, 7)
 Percentage: 16.0%

Test set:
 Shape: (65095, 7)
 Percentage: 20.0%

```
In [5]: # Check set class distribution
print("Train set class distribution:")
print(y_train.value_counts(normalize=True).sort_index())
```

```
Train set class distribution:
navigationalstatus
At anchor                               0.001512
Constrained by her draught              0.037686
Engaged in fishing                      0.015089
Moored                                  0.010355
Power-driven vessel pushing ahead or towing alongside 0.000725
Power-driven vessel towing astern       0.000759
Reserved for future amendment [HSC]    0.005372
Restricted maneuverability              0.004983
Under way sailing                       0.004057
Under way using engine                 0.919462
Name: proportion, dtype: float64
```

```
In [6]: # Check validation set class distribution
print("Validation set class distribution:")
print(y_val.value_counts(normalize=True).sort_index())
```

```
Validation set class distribution:
navigationalstatus
At anchor                               0.001517
Constrained by her draught              0.037696
Engaged in fishing                      0.015094
Moored                                  0.010350
Power-driven vessel pushing ahead or towing alongside 0.000730
Power-driven vessel towing astern       0.000749
Reserved for future amendment [HSC]    0.005358
Restricted maneuverability              0.004974
Under way sailing                       0.004071
Under way using engine                 0.919462
Name: proportion, dtype: float64
```

```
In [7]: # Check test set class distribution
print("Test set class distribution:")
print(y_test.value_counts(normalize=True).sort_index())
```

```
Test set class distribution:
navigationalstatus
At anchor                               0.001505
Constrained by her draught              0.037683
Engaged in fishing                      0.015086
Moored                                  0.010369
Power-driven vessel pushing ahead or towing alongside 0.000722
Power-driven vessel towing astern       0.000753
Reserved for future amendment [HSC]    0.005377
Restricted maneuverability              0.004993
Under way sailing                       0.004056
Under way using engine                 0.919456
Name: proportion, dtype: float64
```

Define Evaluation Helper

Create a function that will provide us a standard to score each model.

```
In [8]: def evaluate_classifier(model, X, y, model_name, split_name, label_encoder=None):
    """
    Evaluate a fitted classifier and return a dictionary of metrics.

    Args:
        model: A fitted sklearn-style classifier with .predict() method
        X: Feature matrix (NumPy array or pandas DataFrame)
        y: True labels
        model_name: String identifier for the model (e.g., "Logistic Regression")
        split_name: String identifier for the data split (e.g., "val" or "test")
        label_encoder: Optional LabelEncoder to decode predictions (required for XGBoost)

    Returns:
        dict: Dictionary containing all computed metrics, predictions, and matrices
    """

    # Get predictions
    y_pred = model.predict(X)

    # Decode predictions if label_encoder provided (for XGBoost)
    if label_encoder is not None:
        y_pred = label_encoder.inverse_transform(y_pred)

    # Compute metrics
    accuracy = accuracy_score(y, y_pred)
    macro_f1 = f1_score(y, y_pred, average="macro")
    weighted_f1 = f1_score(y, y_pred, average="weighted")
    precision_macro = precision_score(y, y_pred, average="macro")
    recall_macro = recall_score(y, y_pred, average="macro")

    # Confusion matrix and classification report
    cm = confusion_matrix(y, y_pred)
    cr = classification_report(y, y_pred)

    # Print concise summary
    print(
        f'{model_name} [{split_name}] | '
        f'Accuracy: {accuracy:.3f} | '
        f'Macro-F1: {macro_f1:.3f} | '
        f'Weighted-F1: {weighted_f1:.3f} | '
        f'Precision (Macro): {precision_macro:.3f} | '
        f'Recall (Macro): {recall_macro:.3f}'
    )

    # Return results dictionary
    return {
        "model_name": model_name,
        "split": split_name,
        "accuracy": accuracy,
        "macro_f1": macro_f1,
        "weighted_f1": weighted_f1,
        "precision_macro": precision_macro,
        "recall_macro": recall_macro,
        "confusion_matrix": cm,
        "classification_report": cr,
    }
```

```
        "y_true": y,
        "y_pred": y_pred
    }

print("Evaluation function instantiated.")
```

Evaluation function instantiated.

Build and Train Models

From `train_models.py`, we are going to use 3 models:

- Logistic Regression:
 - Used as a simple linear baseline that is interpretable and effective for many classification problems.
- Random Forest:
 - Chosen as a robust nonlinear baseline capable of capturing complex patterns and interactions in the data.
- XGBoost:
 - Included as a strong performance baseline due to its efficiency and state-of-the-art results in many classification tasks.

```
In [9]: # Instantiate preprocessor
preprocessor = build_preprocessor()

# Fit preprocessor on training data. Transform splits.
print("Fitting preprocessor on training data...")
X_train_preprocessed = preprocessor.fit_transform(X_train)
X_val_preprocessed = preprocessor.transform(X_val)
X_test_preprocessed = preprocessor.transform(X_test)
print(f"Training data shape after preprocessing: {X_train_preprocessed.shape}")
print(f"Validation data shape after preprocessing: {X_val_preprocessed.shape}")
print(f"Test data shape after preprocessing: {X_test_preprocessed.shape}")
print("Preprocessing complete.\n")

# Build Models
print("Building models...")
logistic_regression_model = build_logistic_regression_model()
random_forest_model = build_random_forest_model()
xgboost_model = build_xgboost()
print("The models are built.\n")

# Train Models
print("Training Logistic Regression...")
logistic_regression_model.fit(X_train_preprocessed, y_train)
print("Logistic Regression training complete.\n")

print("Training Random Forest...")
random_forest_model.fit(X_train_preprocessed, y_train)
print("Random Forest training complete.\n")
```

```
# Encode target labels
print("Prepping target labels for XGBoost...")
label_encoder = LabelEncoder()
y_train_encoded = label_encoder.fit_transform(y_train)
y_val_encoded = label_encoder.transform(y_val)
y_test_encoded = label_encoder.transform(y_test)
print("Target labels encoded.\n")

print("Training XGBoost...")
xgboost_model.fit(X_train_preprocessed, y_train_encoded)
print("XGBoost training complete.\n")
print("\nModel training complete. Awaiting comparison on validation set.")
```

Fitting preprocessor on training data...
 Training data shape after preprocessing: (208300, 23)
 Validation data shape after preprocessing: (52075, 23)
 Test data shape after preprocessing: (65095, 23)
 Preprocessing complete.

Building models...
 The models are built.

Training Logistic Regression...

```
/opt/anaconda3/envs/maritime-ais-ml/lib/python3.11/site-packages/sklearn/linear_model/_logistic.py:1272: FutureWarning: 'multi_class' was deprecated in version 1.5 and will be removed in 1.8. From then on, it will always use 'multinomial'. Leave it to its default value to avoid this warning.
    warnings.warn(
```

Logistic Regression training complete.

Training Random Forest...
 Random Forest training complete.

Prepping target labels for XGBoost...
 Target labels encoded.

Training XGBoost...
 XGBoost training complete.

Model training complete. Awaiting comparison on validation set.

Compare Models on Validation Set

```
In [10]: print("MODEL COMPARISON ON VALIDATION SET\n")

# Evaluate Logistic Regression on validation set
lr_val_results = evaluate_classifier(
    logistic_regression_model,
    X_val_preprocessed,
    y_val,
    "Logistic Regression",
    "val"
)
```

```
# Evaluate Random Forest on validation set
rf_val_results = evaluate_classifier(
    random_forest_model,
    X_val_preprocessed,
    y_val,
    "Random Forest",
    "val"
)

# Evaluate XGBoost (with label encoder) on validation set
xgb_val_results = evaluate_classifier(
    xgboost_model,
    X_val_preprocessed,
    y_val,
    "XGBoost",
    "val",
    label_encoder=label_encoder
)

print("\nVALIDATION SET COMPARISON SUMMARY\n")

# Create a comparison DataFrame
comparison_data = {
    "Model": [
        lr_val_results["model_name"],
        rf_val_results["model_name"],
        xgb_val_results["model_name"]
    ],
    "Accuracy": [
        lr_val_results["accuracy"],
        rf_val_results["accuracy"],
        xgb_val_results["accuracy"]
    ],
    "Macro-F1": [
        lr_val_results["macro_f1"],
        rf_val_results["macro_f1"],
        xgb_val_results["macro_f1"]
    ],
    "Weighted-F1": [
        lr_val_results["weighted_f1"],
        rf_val_results["weighted_f1"],
        xgb_val_results["weighted_f1"]
    ],
    "Precision (Macro)": [
        lr_val_results["precision_macro"],
        rf_val_results["precision_macro"],
        xgb_val_results["precision_macro"]
    ],
    "Recall (Macro)": [
        lr_val_results["recall_macro"],
        rf_val_results["recall_macro"],
        xgb_val_results["recall_macro"]
    ]
}
comparison_df = pd.DataFrame(comparison_data)
```

```
print(comparison_df.to_string(index=False))
print()
```

MODEL COMPARISON ON VALIDATION SET

Logistic Regression [val] | Accuracy: 0.717 | Macro-F1: 0.373 | Weighted-F1: 0.797 | Precision (Macro): 0.303 | Recall (Macro): 0.784
 Random Forest [val] | Accuracy: 0.986 | Macro-F1: 0.875 | Weighted-F1: 0.985 | Precision (Macro): 0.915 | Recall (Macro): 0.844
 XGBoost [val] | Accuracy: 0.978 | Macro-F1: 0.834 | Weighted-F1: 0.978 | Precision (Macro): 0.886 | Recall (Macro): 0.800

VALIDATION SET COMPARISON SUMMARY

| | Model | Accuracy | Macro-F1 | Weighted-F1 | Precision (Macro) | Rec all (Macro) |
|---------------------|----------------------|----------|----------|-------------|-------------------|-----------------|
| Logistic Regression | 0.717216 0.783701 | 0.373483 | | 0.797138 | | 0.303367 |
| Random Forest | 0.985578 0.844384 | 0.875066 | | 0.985190 | | 0.914565 |
| XGBoost | 0.978435 0.799931 | 0.834280 | | 0.977790 | | 0.886171 |

MODEL EVALUATION ON VALIDATION SET

| Model | Accuracy | Macro-F1 | Weighted-F1 | Precision (Macro) | Recall (Macro) |
|---------------------|----------|----------|-------------|-------------------|----------------|
| Logistic Regression | 0.694 | 0.343 | 0.778 | 0.276 | 0.757 |
| Random Forest | 0.985 | 0.881 | 0.984 | 0.934 | 0.845 |
| XGBoost | 0.978 | 0.839 | 0.978 | 0.896 | 0.803 |

Logistic Regression demonstrates noticeably weaker performance on the validation set, particularly in macro-F1 and macro-precision, reflecting limited ability to handle minority classes.

Random Forest stands out as the top performer by a considerable margin; it achieves the highest accuracy and robust scores across all metrics, suggesting strong generalization.

XGBoost also performs well - close to Random Forest - though it trails slightly in all reported metrics, highlighting it as a competitive but marginally less effective alternative on this dataset.

Final Evaluation on Test Set

```
In [11]: print("MODEL EVALUATION ON TEST SET\n")
```

```
# Evaluate Logistic Regression on test set
lr_test_results = evaluate_classifier(
    logistic_regression_model,
    X_test_preprocessed,
    y_test,
    "Logistic Regression",
    "test"
)

# Evaluate Random Forest on test set
rf_test_results = evaluate_classifier(
    random_forest_model,
    X_test_preprocessed,
    y_test,
    "Random Forest",
    "test"
)

# Evaluate XGBoost on test set (with label encoder)
xgb_test_results = evaluate_classifier(
    xgboost_model,
    X_test_preprocessed,
    y_test,
    "XGBoost",
    "test",
    label_encoder=label_encoder
)

print("\nTEST SET COMPARISON SUMMARY\n")

# Create a comparison DataFrame for test set
test_comparison_data = {
    "Model": [
        lr_test_results["model_name"],
        rf_test_results["model_name"],
        xgb_test_results["model_name"]
    ],
    "Accuracy": [
        lr_test_results["accuracy"],
        rf_test_results["accuracy"],
        xgb_test_results["accuracy"]
    ],
    "Macro-F1": [
        lr_test_results["macro_f1"],
        rf_test_results["macro_f1"],
        xgb_test_results["macro_f1"]
    ],
    "Weighted-F1": [
        lr_test_results["weighted_f1"],
        rf_test_results["weighted_f1"],
        xgb_test_results["weighted_f1"]
    ],
    "Precision (Macro)": [
        lr_test_results["precision_macro"],
        rf_test_results["precision_macro"],
        xgb_test_results["precision_macro"]
    ]
}
```

```

        ],
    "Recall (Macro)": [
        lr_test_results["recall_macro"],
        rf_test_results["recall_macro"],
        xgb_test_results["recall_macro"]
    ]
}

test_comparison_df = pd.DataFrame(test_comparison_data)
print(test_comparison_df.to_string(index=False))
print()

```

MODEL EVALUATION ON TEST SET

Logistic Regression [test] | Accuracy: 0.719 | Macro-F1: 0.369 | Weighted-F1: 0.798 | Precision (Macro): 0.300 | Recall (Macro): 0.772
 Random Forest [test] | Accuracy: 0.985 | Macro-F1: 0.883 | Weighted-F1: 0.985 | Precision (Macro): 0.919 | Recall (Macro): 0.860
 XGBoost [test] | Accuracy: 0.979 | Macro-F1: 0.846 | Weighted-F1: 0.978 | Precision (Macro): 0.884 | Recall (Macro): 0.826

TEST SET COMPARISON SUMMARY

| | Model | Accuracy | Macro-F1 | Weighted-F1 | Precision (Macro) | Recall (Macro) |
|---------------------|----------------------|----------|----------|-------------|-------------------|----------------|
| Logistic Regression | 0.718688 0.771674 | 0.369270 | | 0.797992 | | 0.299880 |
| Random Forest | 0.985437 0.859769 | 0.882635 | | 0.984961 | | 0.918861 |
| XGBoost | 0.978647 0.826128 | 0.845623 | | 0.978037 | | 0.884127 |

MODEL EVALUATION ON TEST SET

| Model | Accuracy | Macro-F1 | Weighted-F1 | Precision (Macro) | Recall (Macro) |
|---------------------|----------|----------|-------------|-------------------|----------------|
| Logistic Regression | 0.699 | 0.340 | 0.782 | 0.275 | 0.737 |
| Random Forest | 0.985 | 0.868 | 0.984 | 0.918 | 0.833 |
| XGBoost | 0.979 | 0.842 | 0.978 | 0.896 | 0.807 |

Random Forest remains the best performing model on the test set, achieving the highest accuracy, macro-F1, and macro-precision, showing excellent generalization.

XGBoost also delivers strong results, slightly below Random Forest, while Logistic Regression shows much lower performance, especially on macro-F1 and precision, indicating challenges with the less-represented classes.

Confusion Matrix Visualization for Each Model

```
In [12]: def plot_confusion_matrix(cm, class_names, model_name, normalize=True, save_
    """
        Plot a confusion matrix with proper formatting.

    Args:
        cm: Confusion matrix (numpy array)
        class_names: List of class names
        model_name: Name of the model for title
        normalize: If True, normalize the confusion matrix
        save_path: Path to save the figure (optional)
    """
    if normalize:
        cm_normalized = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        cm_to_plot = cm_normalized
        fmt = '.2f'
        label = 'Normalized'
    else:
        cm_to_plot = cm
        fmt = 'd'
        label = 'Count'

    plt.figure(figsize=(14, 12))
    sns.heatmap(
        cm_to_plot,
        annot=True,
        fmt=fmt,
        cmap='Blues',
        xticklabels=class_names,
        yticklabels=class_names,
        cbar_kws={'label': label}
    )
    plt.title(f'Confusion Matrix - {model_name} (Test Set)', fontsize=16, fontweight='bold')
    plt.ylabel('True Label', fontsize=12)
    plt.xlabel('Predicted Label', fontsize=12)
    plt.xticks(rotation=45, ha='right')
    plt.yticks(rotation=0)
    plt.tight_layout()

    if save_path:
        plt.savefig(save_path, dpi=300, bbox_inches='tight')
        print(f"Saved confusion matrix to: {save_path}")
        plt.show()
    else:
        plt.show()

    class_names = sorted(y_test.unique())

print("Creating confusion matrices for all models on test set...\n")

# Random Forest
```


```
print("Random Forest Confusion Matrix:")
plot_confusion_matrix(
    rf_test_results["confusion_matrix"],
    class_names,
    "Random Forest",
    normalize=True,
    save_path=os.path.join(FIG_DIR, "confusion_matrix_random_forest_test.png")
)

# XGBoost
print("\nXGBoost Confusion Matrix:")
plot_confusion_matrix(
    xgb_test_results["confusion_matrix"],
    class_names,
    "XGBoost",
    normalize=True,
    save_path=os.path.join(FIG_DIR, "confusion_matrix_xgboost_test.png")
)


# Logistic Regression
print("\nLogistic Regression Confusion Matrix:")
plot_confusion_matrix(
    lr_test_results["confusion_matrix"],
    class_names,
    "Logistic Regression",
    normalize=True,
    save_path=os.path.join(FIG_DIR, "confusion_matrix_logistic_regression_te")
```

Creating confusion matrices for all models on test set...

Random Forest Confusion Matrix:
Saved confusion matrix to: ../results/figures/modeling & evaluation/confusion_matrix_random_forest_test.png


**XGBoost Confusion Matrix:**

Saved confusion matrix to: `../results/figures/modeling & evaluation/confusion_matrix_xgboost_test.png`



Logistic Regression Confusion Matrix:

Saved confusion matrix to: ../results/figures/modeling & evaluation/confusion_matrix_logistic_regression_test.png



All three models are most accurate for classes with clear and unique motion patterns, like "Engaged in fishing", "Power-driven vessel towing astern", and "Reserved for future amendment (HSC)". This shows that AIS motion features by themselves are very good at predicting these behaviors.

Logistic Regression performs worse on more subtle classes like At anchor and Moored, which suggests linear methods struggle with complex vessel movements.

Random Forest and XGBoost do better on these tricky classes, but still confuse Moored and At anchor, probably because both tend to have similar slow speeds and low turning rates.

Overall, XGBoost separates almost all classes best, showing that gradient-boosted trees work very well for this type of AIS classification.

Feature Importance on Tree-Based Models

```
In [ ]: def plot_feature_importance(model, feature_names, model_name, top_n=15, save_path=None):
    """
    Plot feature importance as a horizontal bar chart.

    Args:
        model: Trained tree-based model with .feature_importances_ attribute
        feature_names: List of feature names
        model_name: Name of the model for title
        top_n: Number of top features to display
        save_path: Path to save the figure (optional)
    """
    importances = model.feature_importances_

    # Create DataFrame
    importance_df = pd.DataFrame({
        'feature': feature_names,
        'importance': importances
    }).sort_values('importance', ascending=False)

    top_features = importance_df.head(top_n)

    # Plot
    fig, ax = plt.subplots(figsize=(10, max(8, top_n * 0.5)))
    ax.barh(range(len(top_features)), top_features['importance'], color='steelblue')
    ax.set_yticks(range(len(top_features)))
    ax.set_yticklabels(top_features['feature'])
    ax.set_xlabel('Feature Importance', fontsize=12)
    ax.set_title(f'Top {top_n} Feature Importances - {model_name}', fontsize=14)
    ax.invert_yaxis() # Highest importance at top
    plt.tight_layout()

    if save_path:
        plt.savefig(save_path, dpi=300, bbox_inches='tight')
        print(f"Saved feature importance plot to: {save_path}")
        plt_original_show()
        plt.close(fig)
    else:
        plt.show()

    # Feature names from the fitted preprocessor
    print("Extracting feature names from preprocessor...")
    feature_names = preprocessor.get_feature_names_out()
    print(f"Total features after preprocessing: {len(feature_names)}")
    print(f"First 10 features: {feature_names[:10]}")
    print()

    # Random Forest feature importance
    print("Random Forest Feature Importance:")
    plot_feature_importance(
        random_forest_model,
        feature_names,
        "Random Forest",
        top_n=15,
        save_path=os.path.join(FIG_DIR, "feature_importance_random_forest.png"),
    )
```

```
# XGBoost feature importance
print("\nXGBoost Feature Importance:")
plot_feature_importance(
    xgboost_model,
    feature_names,
    "XGBoost",
    top_n=15,
    save_path=os.path.join(FILE_DIR, "feature_importance_xgboost.png"),
)
```

Extracting feature names from preprocessor...


Total features after preprocessing: 23

First 10 features: ['num_sog' 'num_cog' 'num_heading' 'num_width' 'num_length'
 'num_draught' 'cat_shiptype_Cargo' 'cat_shiptype_Dredging'
 'cat_shiptype_Fishing' 'cat_shiptype_HSC']

Random Forest Feature Importance:


Saved feature importance plot to: ../results/figures/modeling & evaluation/feature_importance_random_forest.png

Top 15 Feature Importances - Random Forest



XGBoost Feature Importance:

Saved feature importance plot to: ../results/figures/modeling & evaluation/feature_importance_xgboost.png

Top 15 Feature Importances - XGBoost

```
In [22]: print("TOP FEATURES SUMMARY\n")

rf_importances = pd.DataFrame({
    'feature': feature_names,
    'importance': random_forest_model.feature_importances_
}).sort_values('importance', ascending=False)

xgb_importances = pd.DataFrame({
    'feature': feature_names,
    'importance': xgboost_model.feature_importances_
}).sort_values('importance', ascending=False)

print("\nTop 10 Features - Random Forest:")
print(rf_importances.head(10)[['feature', 'importance']].to_string(index=False))

print("\nTop 10 Features - XGBoost:")
print(xgb_importances.head(10)[['feature', 'importance']].to_string(index=False))
```

TOP FEATURES SUMMARY

Top 10 Features – Random Forest:

| | feature | importance |
|--|----------------------|------------|
| | num_draught | 0.175178 |
| | num_length | 0.162979 |
| | num_sog | 0.161782 |
| | num_width | 0.140289 |
| | cat_shiptype_Tug | 0.065525 |
| | cat_shiptype_Fishing | 0.060011 |
| | num_heading | 0.059539 |
| | num_cog | 0.056410 |
| | cat_shiptype_HSC | 0.048078 |
| | cat_shiptype_Cargo | 0.018688 |

Top 10 Features – XGBoost:

| | feature | importance |
|--|-------------------------------|------------|
| | cat_shiptype_Fishing | 0.295704 |
| | cat_shiptype_HSC | 0.194700 |
| | cat_shiptype_Towing long/wide | 0.072661 |
| | cat_shiptype_Pleasure | 0.069811 |
| | num_draught | 0.058339 |
| | cat_shiptype_Port tender | 0.039720 |
| | cat_shiptype_Towing | 0.034151 |
| | cat_shiptype_Tug | 0.029343 |
| | num_sog | 0.028098 |
| | num_width | 0.027451 |

Random Forest mainly focuses on numeric physical features - especially draught, length, SOG, and width. This suggests that how the vessel moves and its size are very important for predicting navigational status.

On the other hand, XGBoost puts ship type categories (like Fishing, HSC, Towing long/wide) above most numeric features, meaning that what kind of vessel it is also gives strong clues for this model.

Both models rank draught and SOG as top features, showing that vessel load and speed help predict navigational status in any approach.

This shows that Random Forest and XGBoost each use different patterns - the first uses physical movement, the second uses vessel type - to explain navigational status, and both perspectives are useful for predictions.

Save Metrics

```
In [23]: print("SAVING METRICS FOR REPORT...")
val_metrics = pd.DataFrame({
    "Model": [
        lr_val_results["model_name"],
        rf_val_results["model_name"],
        gbm_val_results["model_name"]
    ],
    "Metric": [
        "Accuracy",
        "Precision",
        "Recall",
        "F1 Score"
    ]
})
```

```
        xgb_val_results["model_name"]
    ],
    "Split": [
        lr_val_results["split"],
        rf_val_results["split"],
        xgb_val_results["split"]
    ],
    "Accuracy": [
        lr_val_results["accuracy"],
        rf_val_results["accuracy"],
        xgb_val_results["accuracy"]
    ],
    "Macro_F1": [
        lr_val_results["macro_f1"],
        rf_val_results["macro_f1"],
        xgb_val_results["macro_f1"]
    ],
    "Weighted_F1": [
        lr_val_results["weighted_f1"],
        rf_val_results["weighted_f1"],
        xgb_val_results["weighted_f1"]
    ],
    "Precision_Macro": [
        lr_val_results["precision_macro"],
        rf_val_results["precision_macro"],
        xgb_val_results["precision_macro"]
    ],
    "Recall_Macro": [
        lr_val_results["recall_macro"],
        rf_val_results["recall_macro"],
        xgb_val_results["recall_macro"]
    ]
})
test_metrics = pd.DataFrame({
    "Model": [
        lr_test_results["model_name"],
        rf_test_results["model_name"],
        xgb_test_results["model_name"]
    ],
    "Split": [
        lr_test_results["split"],
        rf_test_results["split"],
        xgb_test_results["split"]
    ],
    "Accuracy": [
        lr_test_results["accuracy"],
        rf_test_results["accuracy"],
        xgb_test_results["accuracy"]
    ],
    "Macro_F1": [
        lr_test_results["macro_f1"],
        rf_test_results["macro_f1"],
        xgb_test_results["macro_f1"]
    ],
    "Weighted_F1": [
        lr_test_results["weighted_f1"],
```

```
        rf_test_results["weighted_f1"],
        xgb_test_results["weighted_f1"]
    ],
    "Precision_Macro": [
        lr_test_results["precision_macro"],
        rf_test_results["precision_macro"],
        xgb_test_results["precision_macro"]
    ],
    "Recall_Macro": [
        lr_test_results["recall_macro"],
        rf_test_results["recall_macro"],
        xgb_test_results["recall_macro"]
    ]
}

# Combine validation and test metrics
all_metrics = pd.concat([val_metrics, test_metrics], ignore_index=True)

# Save metrics to CSV
metrics_path = os.path.join(METRICS_DIR, "model_evaluation_metrics.csv")
all_metrics.to_csv(metrics_path, index=False)
print(f"Saved evaluation metrics to: {metrics_path}")

# Save validation metrics separately
val_metrics_path = os.path.join(METRICS_DIR, "validation_metrics.csv")
val_metrics.to_csv(val_metrics_path, index=False)
print(f"Saved validation metrics to: {val_metrics_path}")

# Save test metrics separately
test_metrics_path = os.path.join(METRICS_DIR, "test_metrics.csv")
test_metrics.to_csv(test_metrics_path, index=False)
print(f"Saved test metrics to: {test_metrics_path}")

# Save confusion matrices (as CSV)
print("\nSaving confusion matrices...")
for model_name, results in [
    ("Logistic_Regression", lr_test_results),
    ("Random_Forest", rf_test_results),
    ("XGBoost", xgb_test_results)
]:
    cm = results["confusion_matrix"]
    cm_df = pd.DataFrame(
        cm,
        index=sorted(y_test.unique()),
        columns=sorted(y_test.unique())
    )
    cm_path = os.path.join(METRICS_DIR, f"confusion_matrix_{model_name}_test")
    cm_df.to_csv(cm_path)
    print(f"Saved confusion matrix to: {cm_path}")

# Save classification reports
print("\nSaving classification reports...")
for model_name, results in [
    ("Logistic_Regression", lr_test_results),
    ("Random_Forest", rf_test_results),
    ("XGBoost", xgb_test_results)
```

```
]:  
    report_path = os.path.join(METRICS_DIR, f"classification_report_{model_r  
with open(report_path, 'w') as f:  
    f.write(f"Classification Report - {model_name} (Test Set)\n")  
    f.write("=" * 80 + "\n\n")  
    f.write(results["classification_report"]))  
    print(f"Saved classification report to: {report_path}")  
  
print("All metrics saved successfully!")
```

SAVING METRICS FOR REPORT...

Saved evaluation metrics to: ../results/metrics/model_evaluation_metrics.csv
Saved validation metrics to: ../results/metrics/validation_metrics.csv
Saved test metrics to: ../results/metrics/test_metrics.csv

Saving confusion matrices...

Saved confusion matrix to: ../results/metrics/confusion_matrix_Logistic_Regr
ession_test.csv
Saved confusion matrix to: ../results/metrics/confusion_matrix_Random_Forest
_test.csv
Saved confusion matrix to: ../results/metrics/confusion_matrix_XGBoost_test.
csv

Saving classification reports...

Saved classification report to: ../results/metrics/classification_report_Log
istic_Regression_test.txt
Saved classification report to: ../results/metrics/classification_report_Ran
dom_Forest_test.txt
Saved classification report to: ../results/metrics/classification_report_XGB
oost_test.txt
All metrics saved successfully!